

Unsupervised Algorithms in machine learning *

Cesar Conejo Villalobos *Data Scientist*

This document provides some examples of unsupervised algorithms in machine learning. In these techniques, we need to infer the properties of the observations without the help of an output variable or *supervisor*. We review two methods: k-means and hierarchical clustering. Then we use some data from Kaggle for applying these techniques to produce a customer segmentation. The platform that we use is R. Because of the number of observations, we are going to use a parallel process for improving the execution times.

Keywords: Unsupervised, algorithms, k-means, hierarchical classification, kaggle, R, parallel

Introduction

In the book *The Elements of Statistical Learning* [Hastie T. \(2008\)](#) explains that in the case of unsupervised learning, data usually has a set of N observations (x_1, \dots, x_N) of a random vector X having joint density $Pr(X)$. The goal is to infer the properties of this probability density.

The techniques that statistics and machine learning offer us for unsupervised learning are the following:

- 1) Principal components, multidimensional scaling.
- 2) Cluster analysis.
- 3) Mixture modeling.
- 4) Association rules.

In this exercise, we are going to focus on cluster analysis. The basis of our model will be the [Kaggle](#) Credit Card dataset for Clustering. The data are an 8950×18 matrix. One variable is categorical and represents the customer ID, the next seventeen are real numbers each representing the behavior of credit cardholders. The goal is to define a marketing strategy based on customer segmentation.

The first aspect we need to solve is to find the number of clusters that we need. [Hastie T. \(2008\)](#) says that we can have two scenarios:

- 1) For data segmentation, the number of clusters is defined as part of the problem, and it is based on the capacity and resources of the company. The goal is to find observations that belong to each proposed group.
- 2) Determine how the observations belong to natural distinct groupings. In this case, the number of clusters is unknown.

For this exercise, we are going to use scenario 2 and trying to find the number of clusters and the characteristics of each group using the following techniques:

- 1) k-means
- 2) Hierarchical clustering.

*Template taken from (<http://github.com/svmiller>). Corresponding author: svmille@clemson.edu.

Techniques

As mentioned before, the goal of unsupervised algorithms is to get the classes as homogeneous as possible and such that they are sufficiently separated. This goal can be specified numerically from the following property:

Suppose that exist a partition $P = (C_1, \dots, C_K)$ of Ω , where g_1, \dots, g_K are the cluster center of the classes:

$$g_k = \frac{1}{|C_k|} \sum_{i \in C_k} x_i$$

Also g is the global center $\frac{1}{N} \sum_{i=1}^N x_i$. We also define:

- Total point scatter: $T = \frac{1}{N} \sum_{i=1}^N \|x_i - g\|^2$.
- Within-cluster point scatter: $W(C) = \frac{1}{N} \sum_{k=1}^K \sum_{i \in C_k} \|x_i - g_k\|^2$.
- Between-cluster point scatter: $B(C) = \sum_{k=1}^K \frac{|C_k|}{N} \|g_k - g\|^2$

In this case, the algorithm requires $B(C)$ to be maximum and $W(C)$ to be minimum. Since the total point scatter T is fixed, then maximizing $B(C)$ automatically minimizes $W(C)$. Therefore, the two goals (homogeneity within classes and separation between classes) are achieved at the same time by minimizing $W(C)$.

Thus, the goal in the *K-means* method is to find a partition C of Ω . Also, we find some representatives of the classes, such that $W(C)$ is minimal. For determining how many clusters a dataset has, we can use the elbow method.

Furthermore, k-means depend on the choice of the number of clusters. On the other hand, hierarchical clustering methods do not expect such designations. Instead, [Hastie T. \(2008\)](#) claims that this method demands the user to specify a measure of dissimilarity between (disjoint) groups of observations, based on the pairwise dissimilarities among the observations.

This method of classification uses a notion of proximity between groups of elements to measure the separation between the classes sought. To do this, the concept of aggregation is introduced, which is nothing more than a dissimilarity between groups of individuals: be $A, B \subseteq \Omega$ then the aggregation between A and B is $\delta(A, B)$. Then we have the following agglomerative clustering methods:

- Single linkage: $\delta_{SL}(A, B) = \min\{d(x_i, d_j) | x_i \in A, x_j \in B\}$
- Complete linkage: $\delta_{CL}(A, B) = \max\{d(x_i, d_j) | x_i \in A, x_j \in B\}$
- Average linkage: $\delta_{AL}(A, B) = \frac{1}{|A||B|} \sum_{x_i \in A, x_j \in B} d(x_i, d_j)$
- Ward linkage: $\delta_{Ward}(A, B) = \frac{|A||B|}{|A|+|B|} \|g_A - g_B\|^2$

Analysis

We are going to define the marketing strategy using k-means and hierarchal clustering. But first, we will see the distribution of the data.

Exploratory Analysis

We create the function `cc_stats()` for analyzing some of the characteristics of the dataset such as:

- Number of complete observations.
- Number of NA values.
- Mean of complete observations.
- Standard deviation of complete observations.
- Number of outliers observations ($Q3 + 1.5IQR$)
- Minimum value of complete observations.
- Maximum value of complete observations.
- 95 quantile
- Upper limit for the value. (mean + 3 sd)

```
---  
# Basic statistics  
# Input: x vector  
# Output: Summary of statistics of the input  
  
cc_stats <- function(x){  
  
  #NA Values  
  nas = sum(is.na(x))  
  
  # Vector with complete values  
  a = x[!is.na(x)]  
  
  # Properties  
  
  m   = mean(a)  
  min = min(a)  
  max = max(a)  
  s   = sd(a)  
  
  # Stats  
  stats <- boxplot.stats(a)  
  n     <- stats$n  
  out   <- length(stats$out)  
  
  Q95 = quantile(a, 0.95)  
  UL  = m + 3*s  
  
  return(c(n     = n,
```

```

    nas    = nas,
    Mean   = m,
    StDev  = s,
    Q_out  = out,
    Min    = min,
    Max    = max,
    Q      = Q95,
    Upper_Limit = UL))
}

```

Using the function *apply()*, we see the statistical characteristics for each of the variables:

```

# Vector with the name of the variable
vars <- c("BALANCE",
          "BALANCE_FREQUENCY",
          "PURCHASES",
          "ONEOFF_PURCHASES",
          "INSTALLMENTS_PURCHASES",
          "CASH_ADVANCE",
          "PURCHASES_FREQUENCY",
          "ONEOFF_PURCHASES_FREQUENCY",
          "PURCHASES_INSTALLMENTS_FREQUENCY",
          "CASH_ADVANCE_FREQUENCY",
          "CASH_ADVANCE_TRX",
          "PURCHASES_TRX",
          "CREDIT_LIMIT",
          "PAYMENTS",
          "MINIMUM_PAYMENTS",
          "PRC_FULL_PAYMENT",
          "TENURE")

```

```

# Apply the function for each variable
describe_stats <- t(data.frame(apply(cc_general[vars], 2, cc_stats)))
describe_stats

```

##	n	nas	Mean	StDev	Q_out	Min	Max
## BALANCE	8950	0	1564.47	2081.53	695	0.000	19043.1
## BALANCE_FREQUENCY	8950	0	0.88	0.24	1493	0.000	1.0
## PURCHASES	8950	0	1003.20	2136.63	808	0.000	49039.6
## ONEOFF_PURCHASES	8950	0	592.44	1659.89	1013	0.000	40761.2
## INSTALLMENTS_PURCHASES	8950	0	411.07	904.34	867	0.000	22500.0
## CASH_ADVANCE	8950	0	978.87	2097.16	1030	0.000	47137.2
## PURCHASES_FREQUENCY	8950	0	0.49	0.40	0	0.000	1.0
## ONEOFF_PURCHASES_FREQUENCY	8950	0	0.20	0.30	782	0.000	1.0
## PURCHASES_INSTALLMENTS_FREQUENCY	8950	0	0.36	0.40	0	0.000	1.0
## CASH_ADVANCE_FREQUENCY	8950	0	0.14	0.20	525	0.000	1.5
## CASH_ADVANCE_TRX	8950	0	3.25	6.82	804	0.000	123.0
## PURCHASES_TRX	8950	0	14.71	24.86	766	0.000	358.0

## CREDIT_LIMIT	8949	1	4494.45	3638.82	248	50.000	30000.0
## PAYMENTS	8950	0	1733.14	2895.06	808	0.000	50721.5
## MINIMUM_PAYMENTS	8637	313	864.21	2372.45	841	0.019	76406.2
## PRC_FULL_PAYMENT	8950	0	0.15	0.29	1474	0.000	1.0
## TENURE	8950	0	11.52	1.34	1366	6.000	12.0
##			Q.95%	Upper_Limit			
## BALANCE	5909.11			7809.07			
## BALANCE_FREQUENCY	1.00			1.59			
## PURCHASES	3998.62			7413.11			
## ONEOFF_PURCHASES	2671.09			5572.10			
## INSTALLMENTS_PURCHASES	1750.09			3124.08			
## CASH_ADVANCE	4647.17			7270.36			
## PURCHASES_FREQUENCY	1.00			1.69			
## ONEOFF_PURCHASES_FREQUENCY	1.00			1.10			
## PURCHASES_INSTALLMENTS_FREQUENCY	1.00			1.56			
## CASH_ADVANCE_FREQUENCY	0.58			0.74			
## CASH_ADVANCE_TRX	15.00			23.72			
## PURCHASES_TRX	57.00			89.28			
## CREDIT_LIMIT	12000.00			15410.90			
## PAYMENTS	6082.09			10418.34			
## MINIMUM_PAYMENTS	2766.56			7981.55			
## PRC_FULL_PAYMENT	1.00			1.03			
## TENURE	12.00			15.53			

First of all, there is only a few values with NA. If we want to see if they both happen at the same time, we can do:

```
sum(is.na(cc_general$CREDIT_LIMIT) & is.na(cc_general$MINIMUM_PAYMENTS))
```

```
## [1] 0
```

As a result, the NA values do not occur in the same row. For fixing these unknown values, we can follow three alternatives:

- Remove the cases.
- Fill in the unknowns using some strategy.
- Use tools that handle these types of values.

In this case, the unknown values only represent 3.51%, so we decide delete that observations.

```
cc_general <- cc_general[-which(is.na(cc_general$CREDIT_LIMIT)
                               | is.na(cc_general$MINIMUM_PAYMENTS)),]
```

Other thing we see is that the variables are measure in diferent scales. For example *BALANCE FREQUENCY*, *PURCHASES FREQUENCY*, *ONE OFF PURCHASES FREQUENCY* and *PURCHASES INSTALLMENTS FREQUENCY* are measure with a score between 0 and 1. Other values are measure in money units and others in number of transactions. Because there are different units then we should scaling variables. We do that with the function *normalize()*:

```

---
## Normalize
## Input: Numeric vector
## Output: Vector normalized.

normalize <- function(x){

  min_x <- min(x)
  max_x <- max(x)

  return((x - min_x)/(max_x - min_x))

}

```

Then, we apply the function to each variable

```
cc_general_norm <- data.frame(apply(cc_general[vars], 2, normalize))
```

Finally, we apply `cc_stats()` again for seeing the changes in our data.

```
describe_stats_norm <- t(data.frame(apply(cc_general_norm[vars], 2, cc_stats)))
describe_stats_norm
```

	n	nas	Mean	StDev	Q_out	Min	Max	Q.95%
## BALANCE	8636	0	0.084	0.110	666	0	1	0.312
## BALANCE_FREQUENCY	8636	0	0.895	0.208	1511	0	1	1.000
## PURCHASES	8636	0	0.021	0.044	767	0	1	0.083
## ONEOFF_PURCHASES	8636	0	0.015	0.041	961	0	1	0.067
## INSTALLMENTS_PURCHASES	8636	0	0.019	0.041	811	0	1	0.080
## CASH_ADVANCE	8636	0	0.021	0.045	976	0	1	0.100
## PURCHASES_FREQUENCY	8636	0	0.496	0.401	0	0	1	1.000
## ONEOFF_PURCHASES_FREQUENCY	8636	0	0.206	0.300	749	0	1	1.000
## PURCHASES_INSTALLMENTS_FREQUENCY	8636	0	0.369	0.398	0	0	1	1.000
## CASH_ADVANCE_FREQUENCY	8636	0	0.092	0.135	346	0	1	0.389
## CASH_ADVANCE_TRX	8636	0	0.027	0.056	794	0	1	0.122
## PURCHASES_TRX	8636	0	0.042	0.070	716	0	1	0.165
## CREDIT_LIMIT	8636	0	0.149	0.122	243	0	1	0.399
## PAYMENTS	8636	0	0.035	0.057	785	0	1	0.121
## MINIMUM_PAYMENTS	8636	0	0.011	0.031	841	0	1	0.036
## PRC_FULL_PAYMENT	8636	0	0.159	0.296	1343	0	1	1.000
## TENURE	8636	0	0.922	0.218	1290	0	1	1.000
##	Upper_Limit							
## BALANCE	0.41							
## BALANCE_FREQUENCY	1.52							
## PURCHASES	0.15							
## ONEOFF_PURCHASES	0.14							
## INSTALLMENTS_PURCHASES	0.14							
## CASH_ADVANCE	0.16							

## PURCHASES_FREQUENCY	1.70
## ONEOFF_PURCHASES_FREQUENCY	1.11
## PURCHASES_INSTALLMENTS_FREQUENCY	1.56
## CASH_ADVANCE_FREQUENCY	0.50
## CASH_ADVANCE_TRX	0.20
## PURCHASES_TRX	0.25
## CREDIT_LIMIT	0.52
## PAYMENTS	0.21
## MINIMUM_PAYMENTS	0.10
## PRC_FULL_PAYMENT	1.05
## TENURE	1.58

Now, all the variables are in a scale from 0 to 1. Also, there is no change in the variance of the variables.

K-means

For applying the *k-means*, we develop the following code. We set the seed 1234 for reproducibility purposes.

```
set.seed(1234)
```

First, we need decide the number of cluster. We apply K-means clustering to the data using the following techniques:

- Hartigan-Wong
- MacQueen
- Lloyd
- Forgy

```
library(snow)

# How many k?

cl <- makeCluster(4, type="SOCK")

ignore <- clusterEvalQ(cl, {library(MASS); NULL})

#Hartigan-Wong
results_HW <- lapply(seq(1,20), function(x) kmeans(cc_general_norm,
                                                    centers = x,
                                                    algorithm = "Hartigan-Wong",
                                                    nstart = 20))

variance_HW <- sapply(results_HW, function(results_HW) results_HW$tot.withinss)
```

```

# MacQueen
results_MQ <- lapply(seq(1,20), function(x) kmeans(cc_general_norm,
                                                    centers = x,
                                                    algorithm = "MacQueen",
                                                    nstart = 20))

variance_MQ <- sapply(results_MQ, function(results_MQ) results_MQ$tot.withinss)

# Lloyd
results_Ll <- lapply(seq(1,20), function(x) kmeans(cc_general_norm,
                                                    centers = x,
                                                    algorithm = "Lloyd",
                                                    nstart = 20))

variance_Ll <- sapply(results_Ll, function(results_Ll) results_Ll$tot.withinss)

# Forgy
results_Fg <- lapply(seq(1,20), function(x) kmeans(cc_general_norm,
                                                    centers = x,
                                                    algorithm = "Forgy",
                                                    nstart = 20))

variance_Fg <- sapply(results_Fg, function(results_Fg) results_Fg$tot.withinss)

stopCluster(cl)

```


Then, we can observe the total within-cluster sum of squares for K-means clustering for some cluster from 1 to 20.

```

plot(variance_HW,
     col = "red",
     type = "b",
     xlab = "Number of cluster k",
     ylab = "Sum of squares",
     main = "Elbow Method: No. of clusters by algorithm")
points(variance_MQ, col = "blue", type = "b")
points(variance_Ll, col = "green", type = "b")
points(variance_Fg, col = "magenta", type = "b")
legend("topright",
     legend = c("Hartigan", "MacQueen", "Lloyd", "Forgy"),
     col = c("red", "blue", "green", "magenta"),
     lty = 1,
     lwd = 1)

```

We can see that the kink occur at $k = 4$, so this is the number of cluster that we propose for the



marketing strategy.

Now, the question that we need to answer is which method for k-means to use. Therefore, we code the following lines that shows the results of the clustering in 4 groups.

```
# Which method:

# Hartigan-Wong
results_HW <- lapply(seq(1,20), function(x) kmeans(cc_general_norm,
                                                    centers = 4,
                                                    nstart = 20,
                                                    algorithm = "Hartigan-Wong"))
betweeness_HW <- sapply(results_HW, function(results_HW) results_HW$betweeness)

# MacQueen
results_MQ <- lapply(seq(1,20), function(x) kmeans(cc_general_norm,
                                                    centers = 4,
                                                    nstart = 20,
                                                    algorithm = "MacQueen"))
betweeness_MQ <- sapply(results_MQ, function(results_MQ) results_MQ$betweeness)

# Lloyd
results_Ll <- lapply(seq(1,20), function(x) kmeans(cc_general_norm,
                                                    centers = 4,
                                                    nstart = 20,
```

```

                                algorithm = "Lloyd"))
betweeness_Ll <- sapply(results_Ll, function(results_Ll) results_Ll$betweeness)

# Forgy
results_Fg <- lapply(seq(1,20), function(x) kmeans(cc_general_norm,
                                centers = 4,
                                nstart = 20,
                                algorithm = "Forgy"))
betweeness_Fg <- sapply(results_Fg, function(results_Fg) results_Fg$betweeness)

AVG_Between_Class <- data.frame(Method = c("Hartigan-Wong", "MacQueen",
                                "Lloyd", "Forgy"),
                                AVG_Between_Class = c(mean(betweeness_HW), mean(betweeness_MQ),
                                mean(betweeness_Ll), mean(betweeness_Fg))
)

```

Method	Avg between
Hartigan-Wong	3222.460
MacQueen	3221.499
Lloyd	3220.927
Forgy	3221.112

Because Lloyd reaches the minimum average between-cluster point scatter, we choose that method. As a result, we code the method in the following way:

```

Cluster_FG <- kmeans(x = cc_general_norm,
                    centers = 4,
                    iter.max = 100,
                    nstart = 20,
                    algorithm = "Lloyd")

```

Finally, we can see some properties of the cluster. For example, the number of observations by cluster is:

```
Cluster_FG$size
```

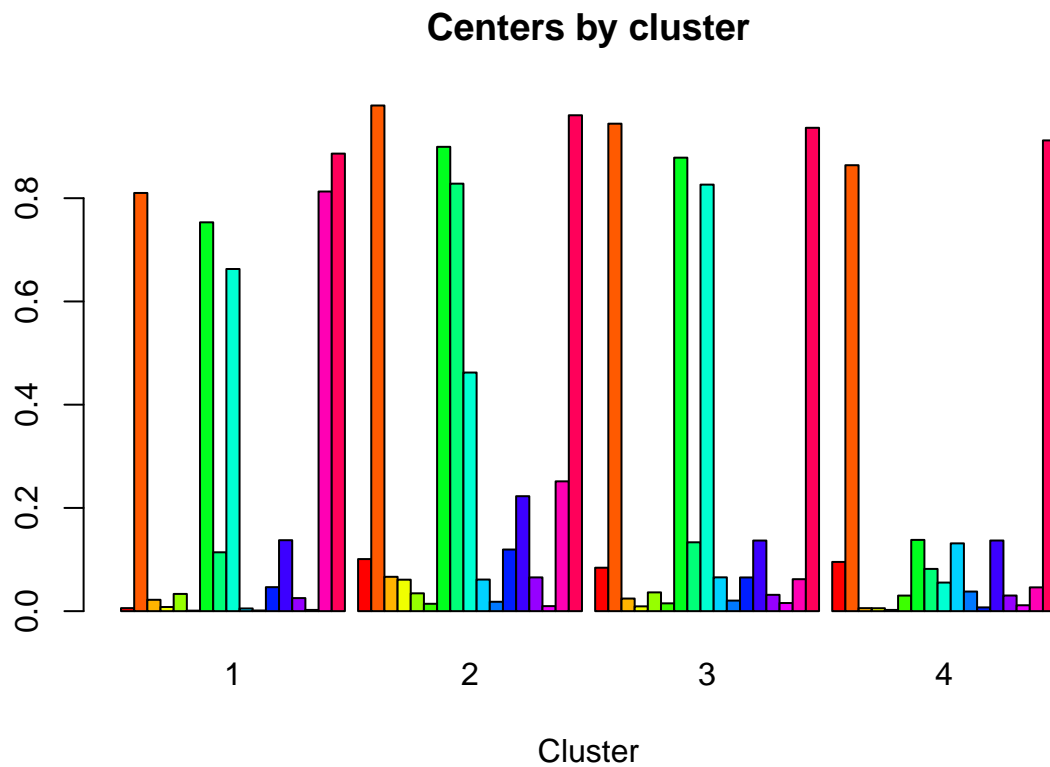
```
## [1] 896 1249 2146 4345
```

In this case, we can see that the classes are well balanced. Also, we can see the center of each cluster and realize some interpretations.

```

barplot(t(Cluster_FG$centers),
        main = "Centers by cluster",
        xlab = "Cluster",
        beside = TRUE,
        col = rainbow(17)
)

```



Hierachical cluster


```

model_Ward <- hclust(dist(cc_general_norm), method = "ward.D2")

plot(model_Ward, labels = FALSE)
rect.hclust(model_Ward, k = 4, border = "blue")

```

Cluster Dendrogram



```
dist(cc_general_norm)
hclust (*, "ward.D2")
```

```
cluster <- cutree(model_Ward, k = 4)


cc_general_norm_cluster <- cbind(cc_general_norm, cluster)

library(rattle)

cc_general_center <- centers.hclust(cc_general_norm,
                                   model_Ward,
                                   nclust = 4,
                                   use.median = FALSE)
```

```
barplot(t(cc_general_center),
        beside = TRUE,
        main = "Cluster Interpretation",
        col = rainbow(17)
)
```

Cluster Interpretation



```
## radar graph
center <- as.data.frame(cc_general_center)
maximos <- apply(center,2,max)
minimos <- apply(center,2,min)
center <- rbind(minimos,center)
center <- rbind(maximos,center)
```

```
library(fmsb)
radarchart(center,
  maxmin = TRUE,
  axistype = 4,
  axislabcol = "slategray4",
  centerzero = FALSE,
  seg = 8,
  cglcol = "gray67",
  pcol=c("green","blue","red", "purple", "black", "brown"),
  plty = 1,
  plwd = 5,
  title = "Comparación de clústeres")
```

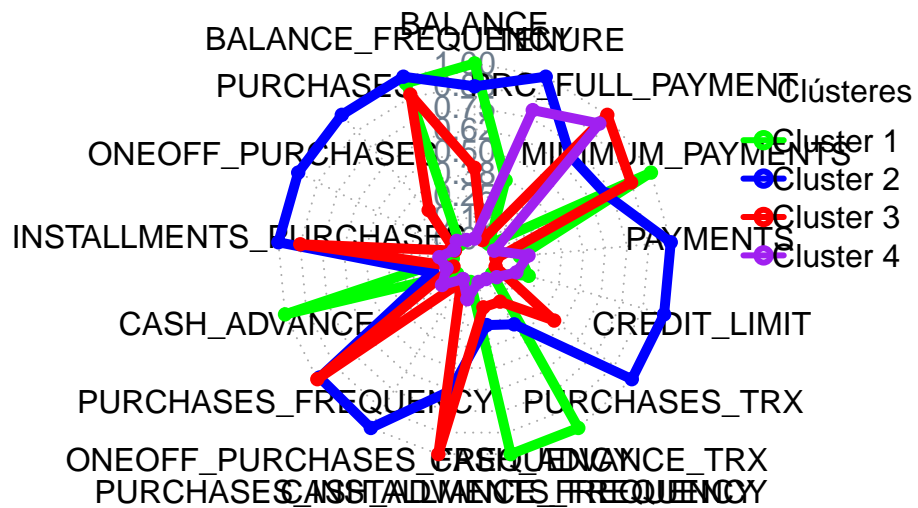
```
legenda <- legend(1.5,1, legend=c("Cluster 1","Cluster 2","Cluster 3", "Cluster 4"),
```

```

seg.len=-1.4,
title="Clústeres",
pch=21,
bty="n" ,lwd=3, y.intersp=1, horiz=FALSE,
col=c("green","blue","red", "purple", "black", "brown")
)

```

Comparación de clústeres



Conclusion

References

Hastie T., Tibshirani R., Friedman J. 2008. *The elements of Statistical Learning*. Springer.